



SUMMIT
ONLINE

TAIWAN | SEPTEMBER 02, 2021

KEY 402

Amazon Networking Foundations

Shih-Yong Wang

Senior Solutions Architect
Amazon Web Services





AWS recognized as a Cloud Leader for the 10th consecutive year ¹

Gartner, Magic Quadrant for Cloud Infrastructure as a Service, Worldwide, Raj Bala, Bob Gill, Dennis Smith, David Wright, July 2019. ID G00365830. Gartner does not endorse any vendor, product or service depicted in its research publications, and does not advise technology users to select only those vendors with the highest ratings. Gartner research publications consist of the opinions of Gartner's research organization and should not be construed as statements of fact. Gartner disclaims all warranties, expressed or implied, with respect to this research, including any warranties of merchantability or fitness for a particular purpose. The Gartner logo is a trademark and service mark of Gartner, Inc., and/or its affiliates, and is used herein with permission. All rights reserved.

Figure 1. Magic Quadrant for Cloud Infrastructure and Platform Services

¹ <https://aws.amazon.com/blogs/aws/aws-named-as-a-cloud-leader-for-the-10th-consecutive-year-in-gartners-infrastructure-platform-services-magic-quadrant/>

AWS Global Infrastructure

AWS global infrastructure

25 geographical regions, 81 availability zones, 230+ POPs

Region & Number of Availability Zones (AZs)

GovCloud (US)

US-East (3), US-West (3)

US West

Oregon (4)

Northern California (3)

US East

N. Virginia (6), Ohio (3)

Canada

Central (3)

South America

São Paulo (3)

Africa

Cape Town (3)

Europe

Frankfurt (3), Paris (3),
Ireland (3), Stockholm (3),
London (3), Milan (3)

Middle East

Bahrain (3)

Asia Pacific

Singapore (3), Sydney (3),
Tokyo (4), Osaka (3)

Seoul (4), Mumbai (3),
Hong Kong (3)

China

Beijing (3), Ningxia (3)



Announced Regions

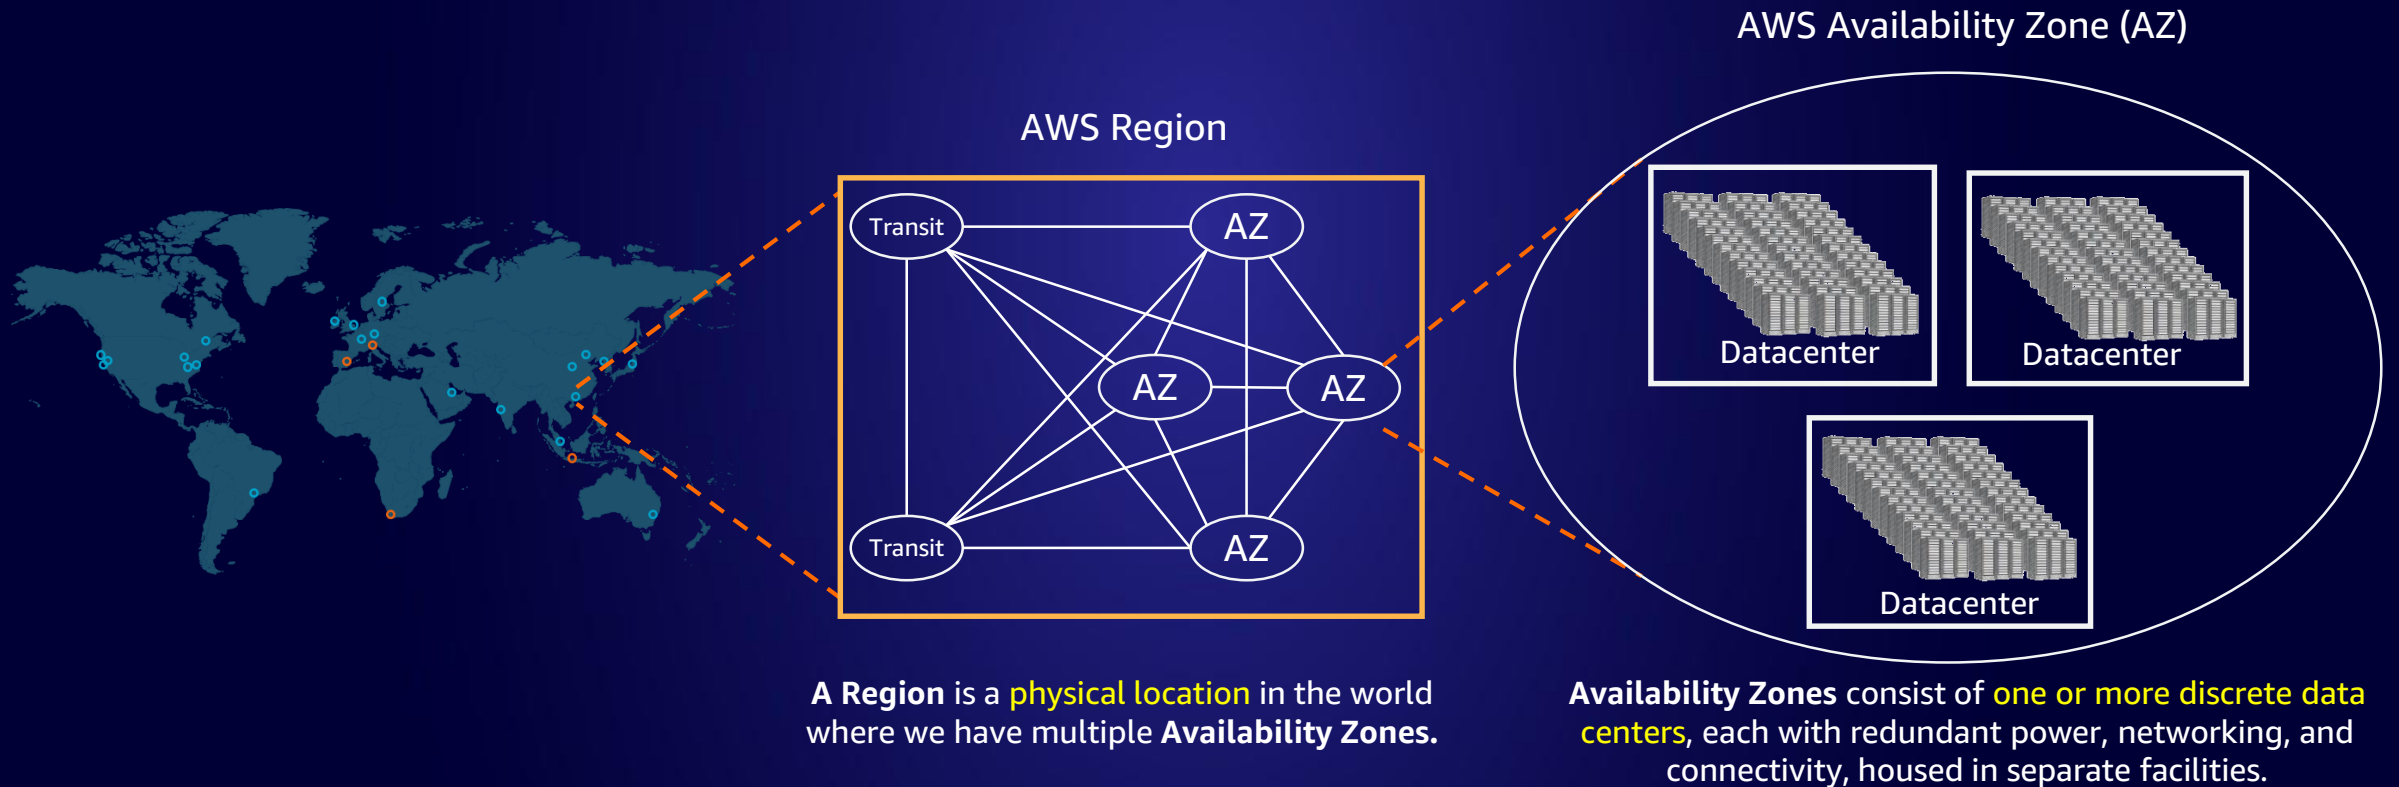
7 Regions and 21 AZs in Australia, India, Indonesia, Israel, Spain, Switzerland, and United Arab Emirates (UAE)

<https://aws.amazon.com/about-aws/global-infrastructure/>



AWS Region design

AWS Regions are comprised of multiple AZs for **high availability**, **high scalability**, and **high fault tolerance**. Applications and data are replicated in real time and consistent in the different AZs.



AWS Availability Zone (AZ) design

- **Fully isolated** infrastructure with one or more datacenters
- **Meaningful distance** of separation
- **Unique power** infrastructure
- Many 100Ks of **servers at scale**
- Datacenters connected via fully redundant and isolated metro **fiber**




Intra-AZ & inter-AZ connectivity

- **Dark fiber** spans
 - Optimized for low-latency & physical diversity
 - High fiber density
 - Amazon controlled infrastructure
- Dense wavelength division multiplexing (**DWDM**) for inter-AZ connections
- **Optical level failover** reducing the impact of physical faults



AWS network connection design




Amazon VPC


Amazon Virtual Private Cloud (Amazon VPC)

- Lets you provision a **logically isolated** section of the AWS Cloud
- You can launch **AWS resources** in a virtual network that you define
- You have **complete control** over your virtual networking environment, including selection of your own IP address range, creation of subnets, and configuration of route tables and network gateways
- You can use both **IPv4 and IPv6** in your VPC

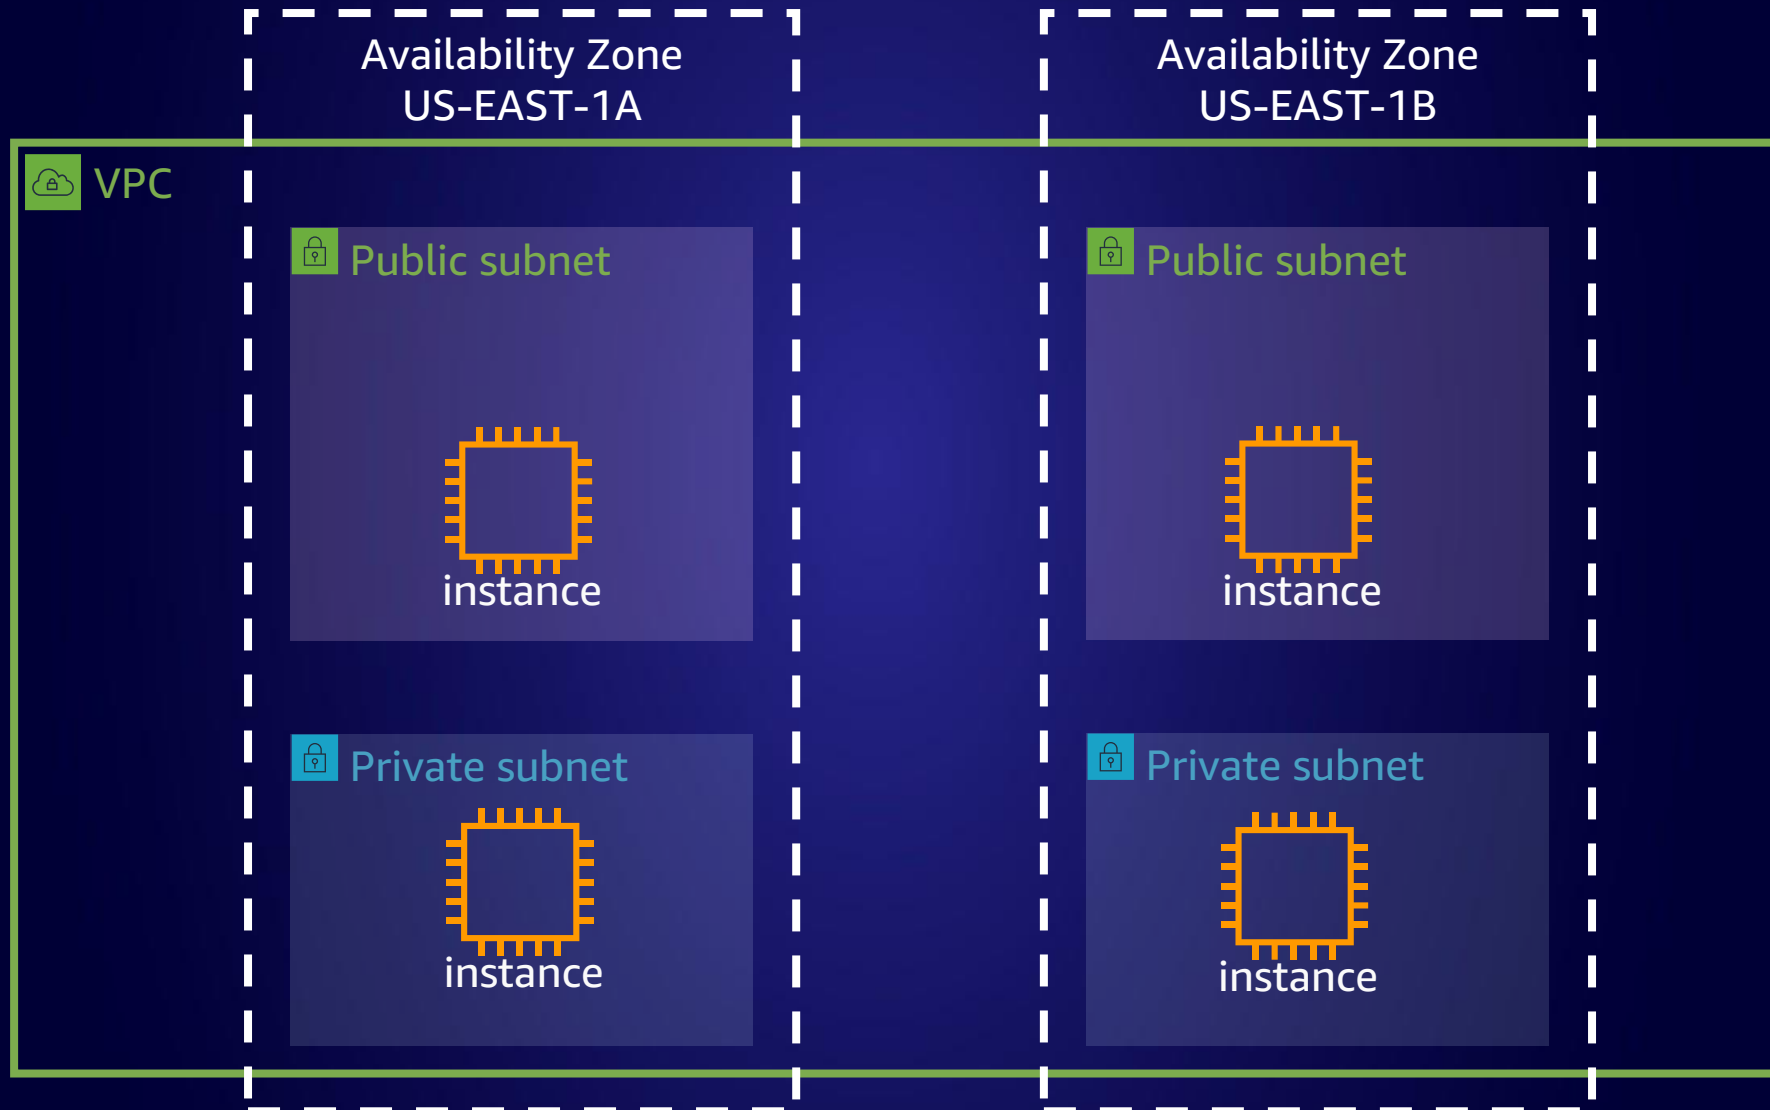
Amazon Virtual Private Cloud (VPC)



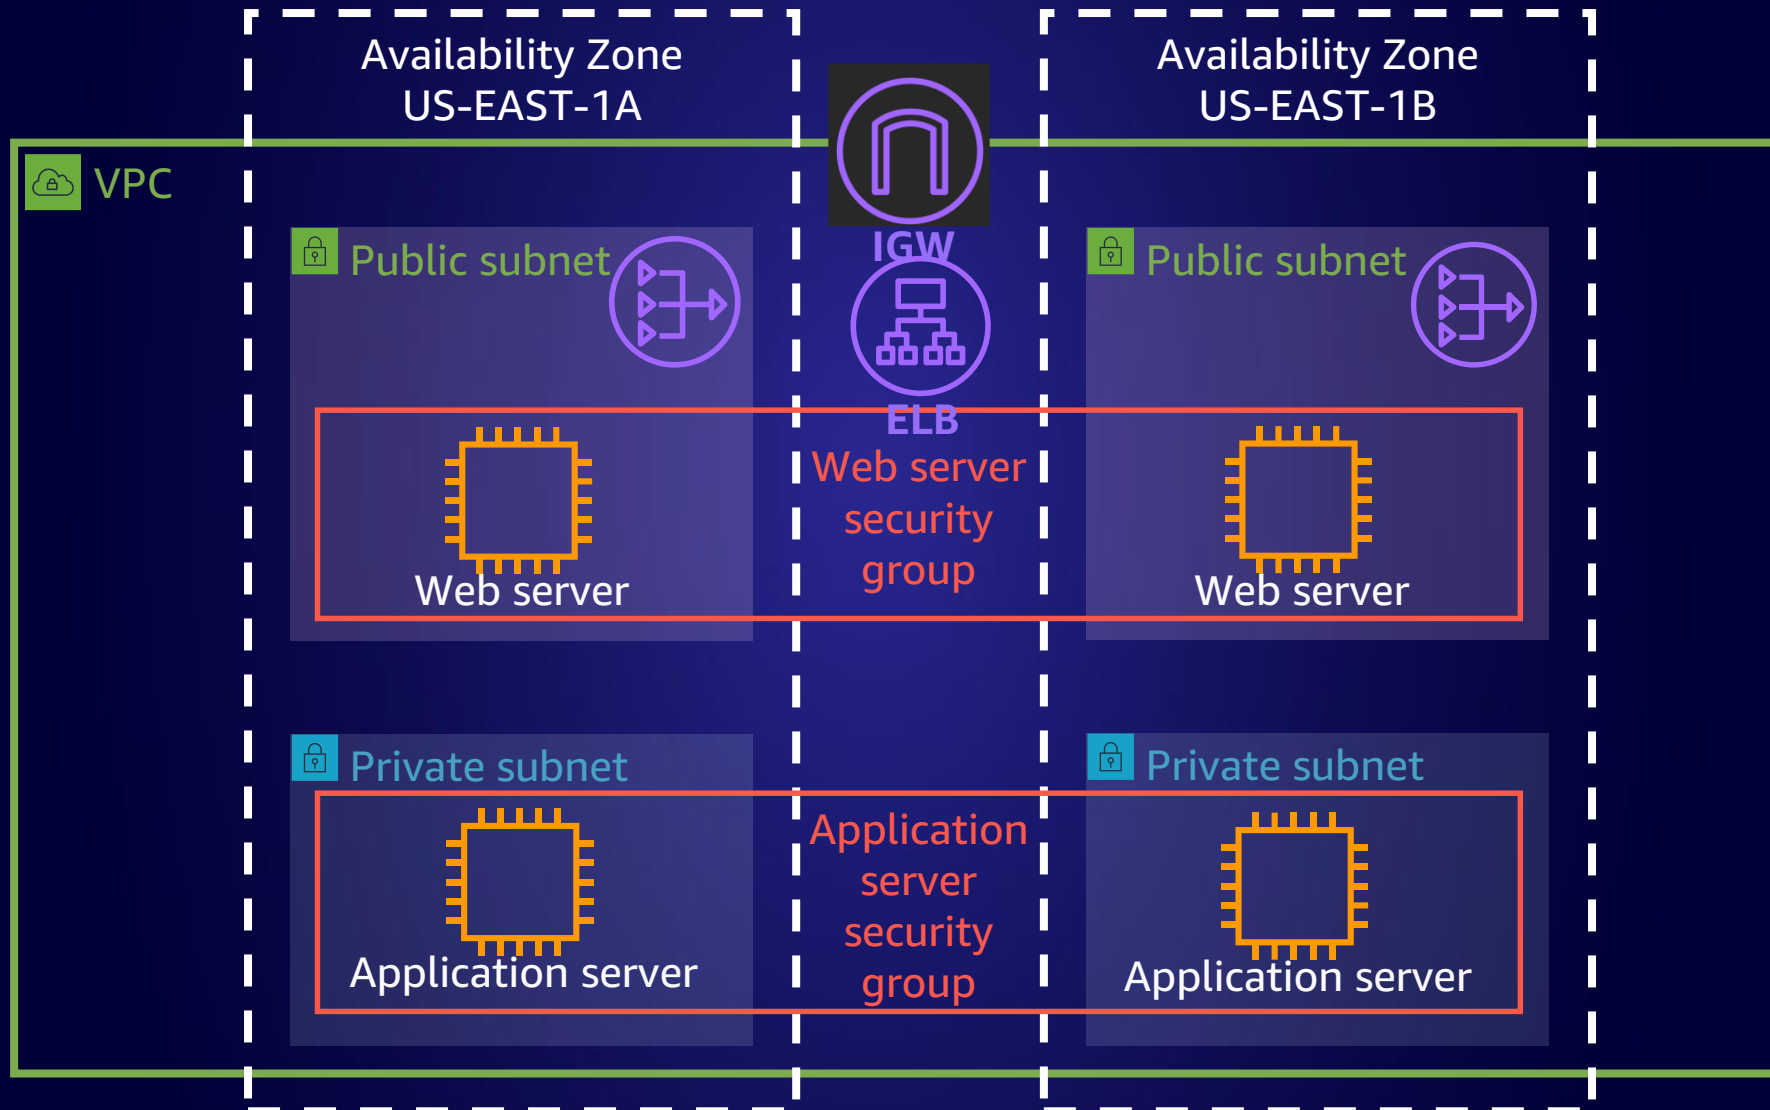
Subnets



Elastic Compute Cloud (EC2) instances



Example web application




Amazon Regions and AZ

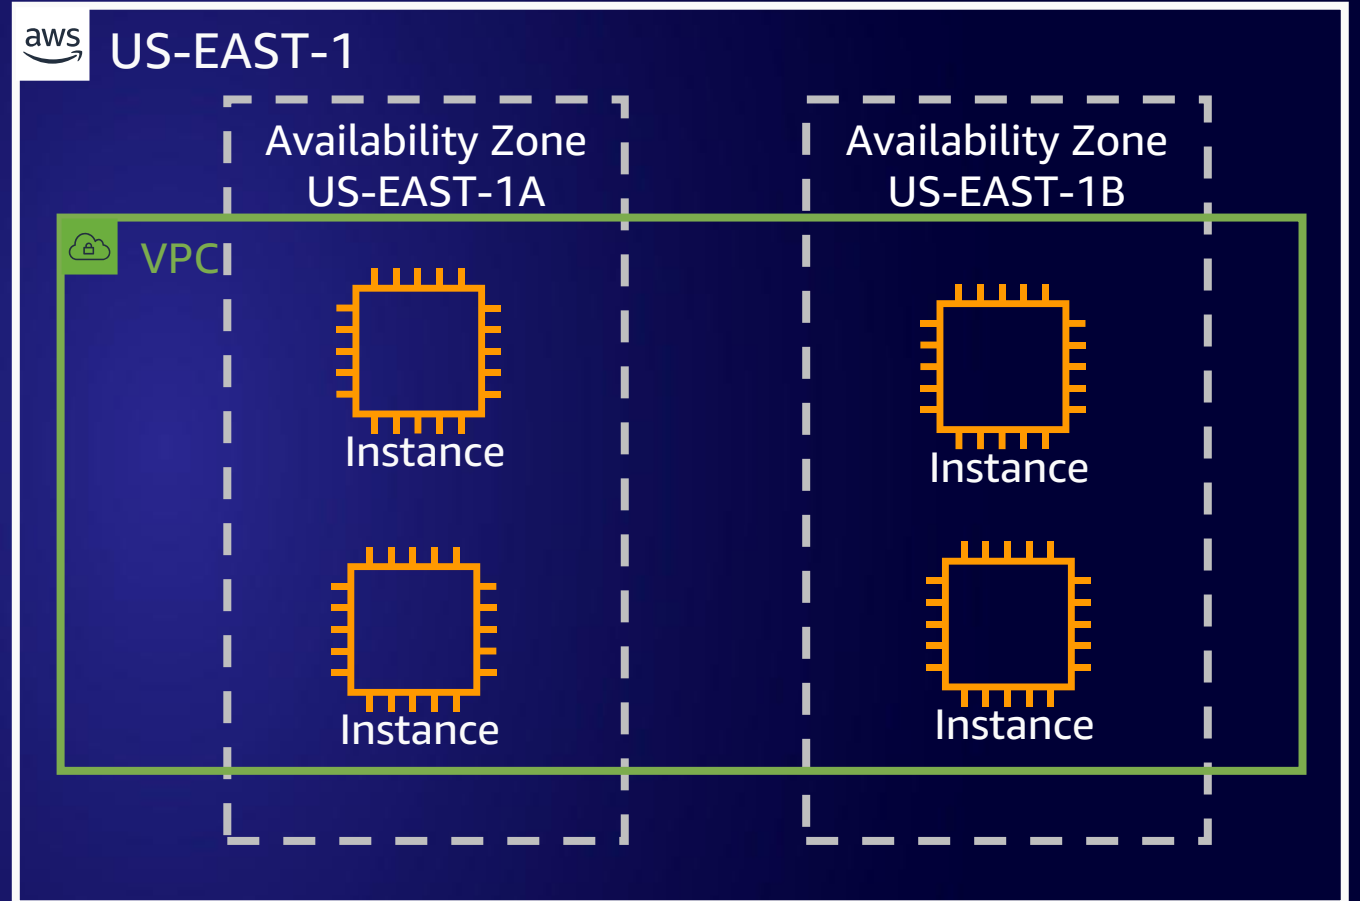
Region



Availability Zone



Data center, rack, host



Gateways

Amazon Internet Gateway (IGW)

- A horizontally scaled, redundant, and highly available VPC **component**
- Allows **communication** between instances in your VPC and the internet
- No availability **risks or bandwidth** constraints on your network traffic

Amazon NAT Gateway (NAT GW)


- In a **private** subnet
 - Instance **can** connect to services outside your VPC
 - But external services **cannot** initiate a connection with those instances
- When sending response traffic to the instances, the NAT device **translates** the addresses back to the original source IPv4 addresses
- Each NAT gateway is created in a specific **Availability Zone** and implemented with redundancy in that zone

Internet Gateway and NAT Gateway



Connecting to Other VPCs

VPC peering



- Full **private IP connectivity** between two VPCs
- Can peer VPCs **across regions**
- VPCs can be in **different accounts**
- VPC CIDR ranges must **not** overlap

VPC peering: Things to know

- **Can** reference security groups from the peer VPC in the same Region
- **Can** enable DNS hostname resolution to return private IP addresses
- **Can** peer for both IPv4 and IPv6 addresses
- **Cannot** have overlapping IP addresses
- **Cannot** have multiple peers between the same pair of VPCs
- **Cannot** use jumbo frames across inter-region VPC peering

The calculation sample 1

Full mesh: How many Amazon VPC peering connections do I need (full mesh)?

VPC x 10

45

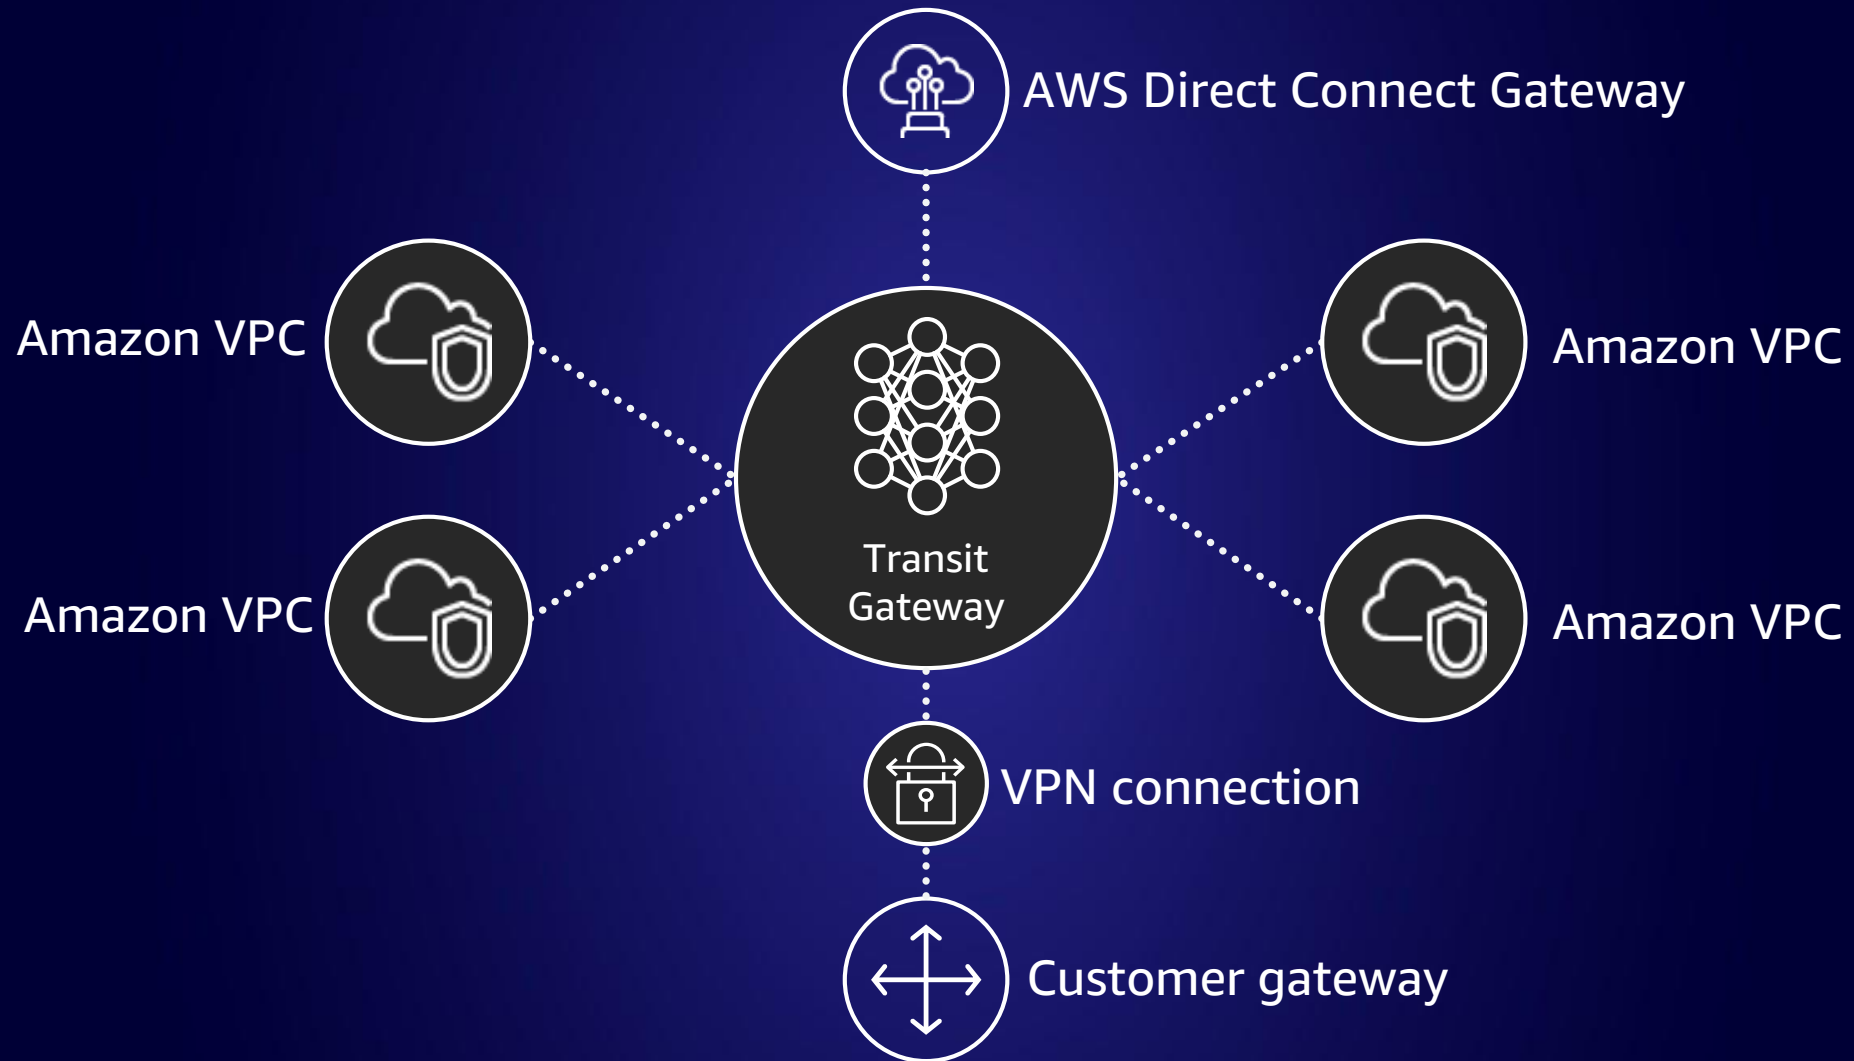
The calculation sample 2

Full mesh: How many Amazon VPC peering connections do I need (full mesh)?

VPC x 100

4,500

With AWS Transit Gateway ...



AWS Transit Gateway capability (1/2)

TGWs per account/TGW
attachments per Amazon VPC

5

Maximum burstable
bandwidth per attachment

50 Gbps

AWS Transit Gateway capability (2/2)

Routes per TGW


10,000

Number of TGW
attachments per region
per account

5,000

Connectivity to on-premises networks

AWS site-to-site VPN

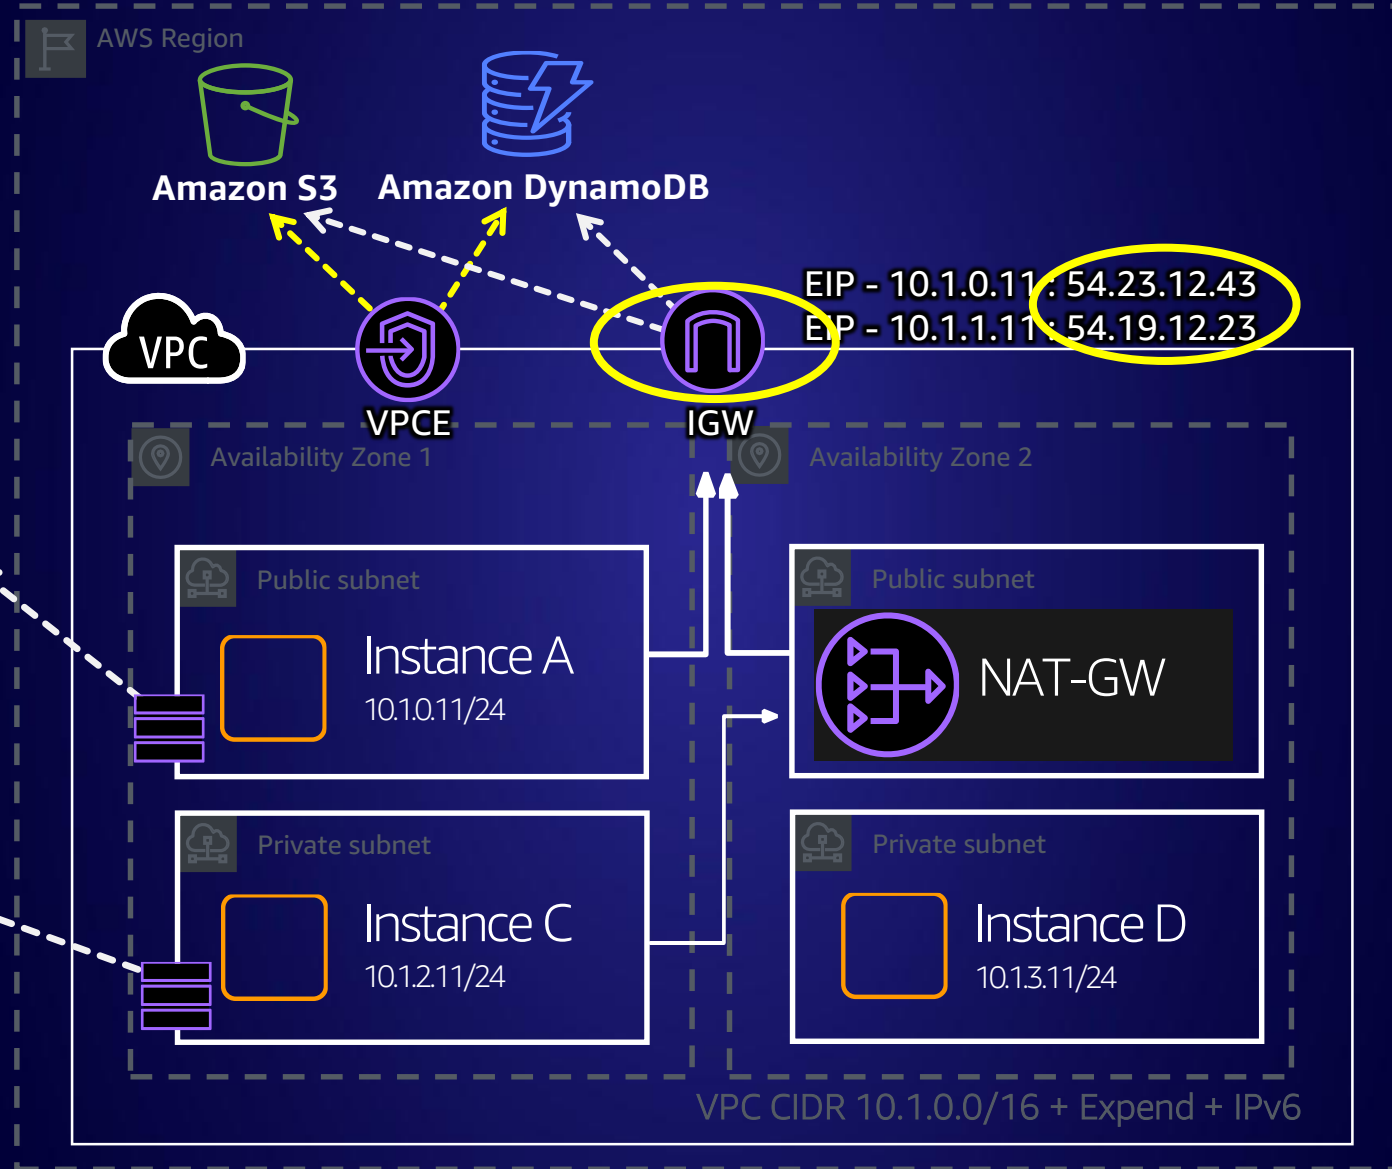


1x VPN connection = 2x VPN tunnels

1x VPN tunnel = 1.25 Gbps

Amazon VPC endpoints

Amazon VPC Gateway endpoints



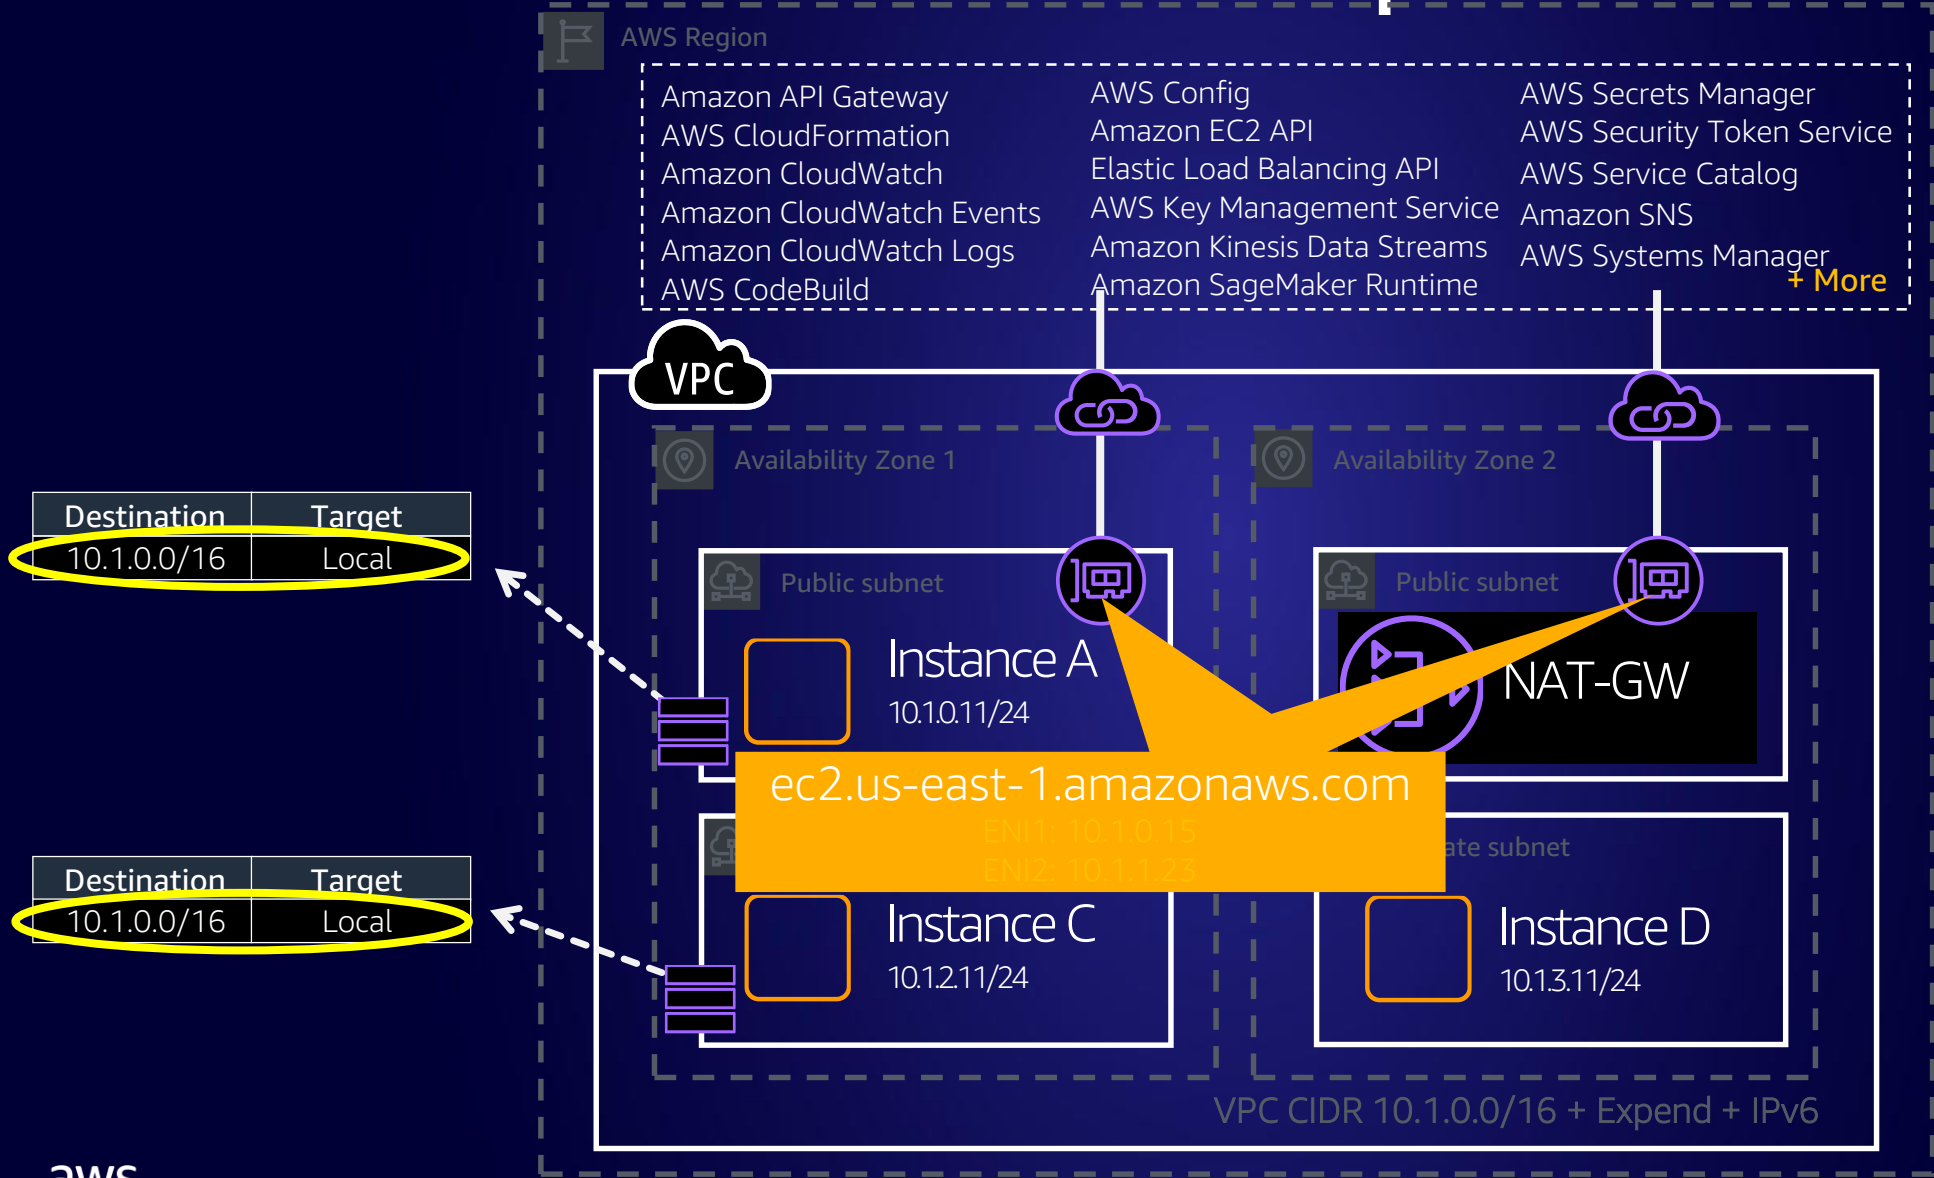
The
internet

VPCE =
Virtual Private Endpoint
(Type: Gateway)

Destination	Target
10.1.0.0/16	Local
0.0.0.0/0	IGW
S3.prefix.list	VPCE-123

Destination	Target
10.1.0.0/16	Local
0.0.0.0/0	NAT-GW
DDB.prefix.list	VPCE-123

Amazon VPC Interface endpoints



Tens of services now supported over AWS PrivateLink

AWS PrivateLink can reach public services, privately from your VPC

No routes needed! (almost)

Amazon VPC endpoints console

Endpoints > Create Endpoint

Create Endpoint

A VPC endpoint allows you to securely connect your VPC to another service. An interface endpoint is an elastic network interface (ENI) that serves as an entry point for traffic destined to the service. A gateway endpoint serves as a target for a route in your route table for traffic destined for the service.

Service Name Select a service ⓘ

<input type="radio"/>	com.amazonaws.us-east-2.dynamodb	amazon	Gateway
<input type="radio"/>	com.amazonaws.us-east-2.ec2	amazon	Interface
<input type="radio"/>	com.amazonaws.us-east-2.ec2messages	amazon	Interface
<input type="radio"/>	com.amazonaws.us-east-2.elasticloadbala...	amazon	Interface
<input type="radio"/>	com.amazonaws.us-east-2.kinesis-streams	amazon	Interface
<input type="radio"/>	com.amazonaws.us-east-2.s3	amazon	Gateway
<input type="radio"/>	com.amazonaws.us-east-2.servicecatalog	amazon	Interface
<input type="radio"/>	com.amazonaws.us-east-2.ssm	amazon	Interface

Type: Gateway

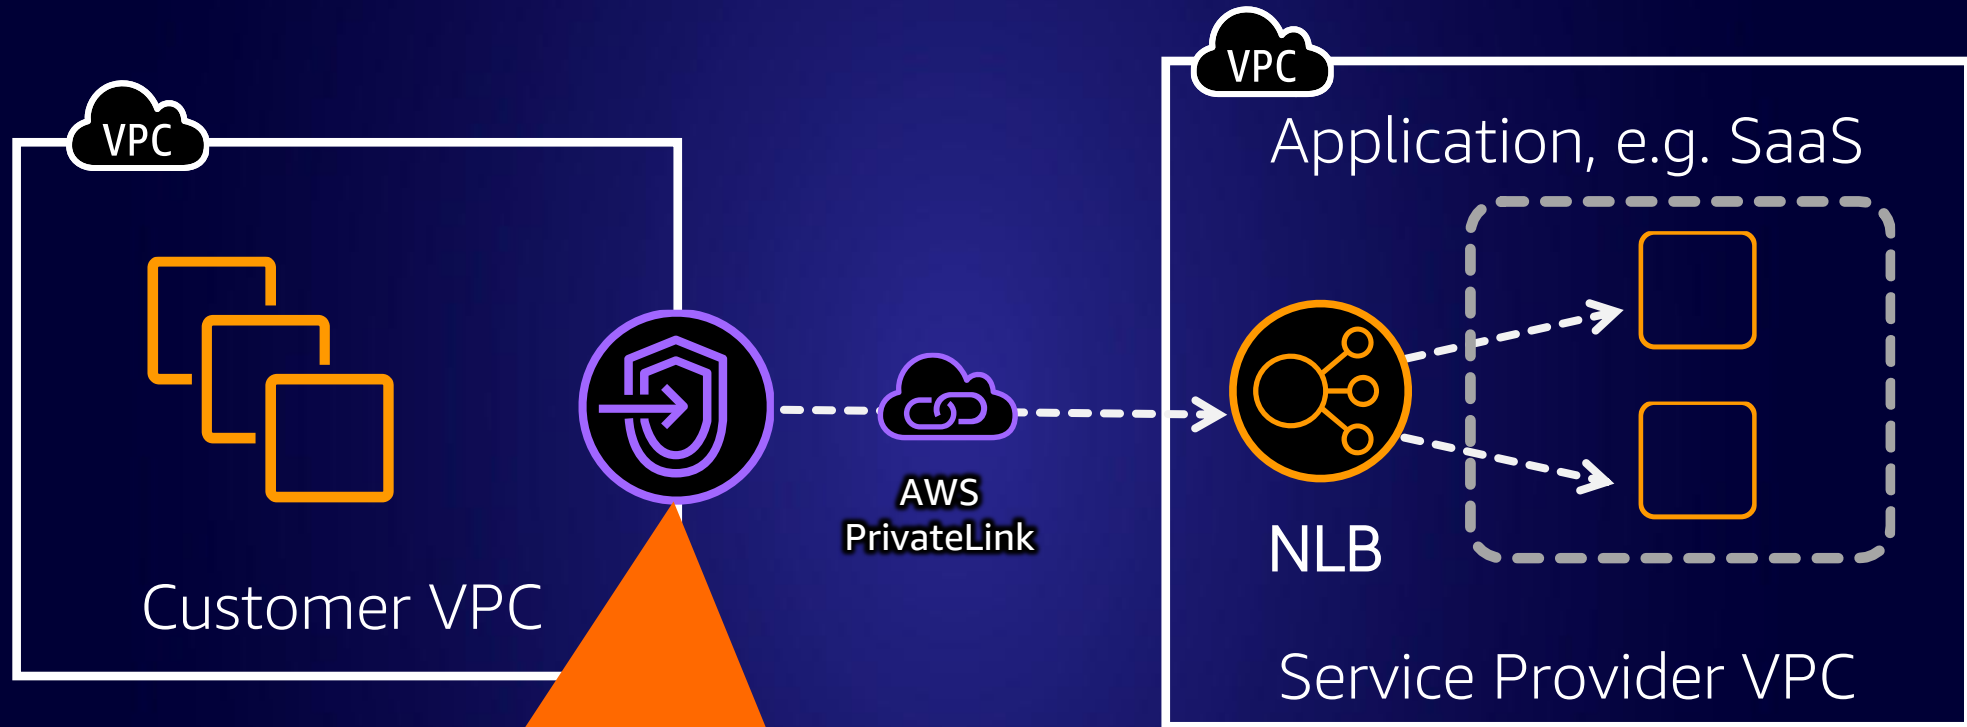
Type: Interface

* Required

Cancel

Create endpoint


AWS PrivateLink for service providers



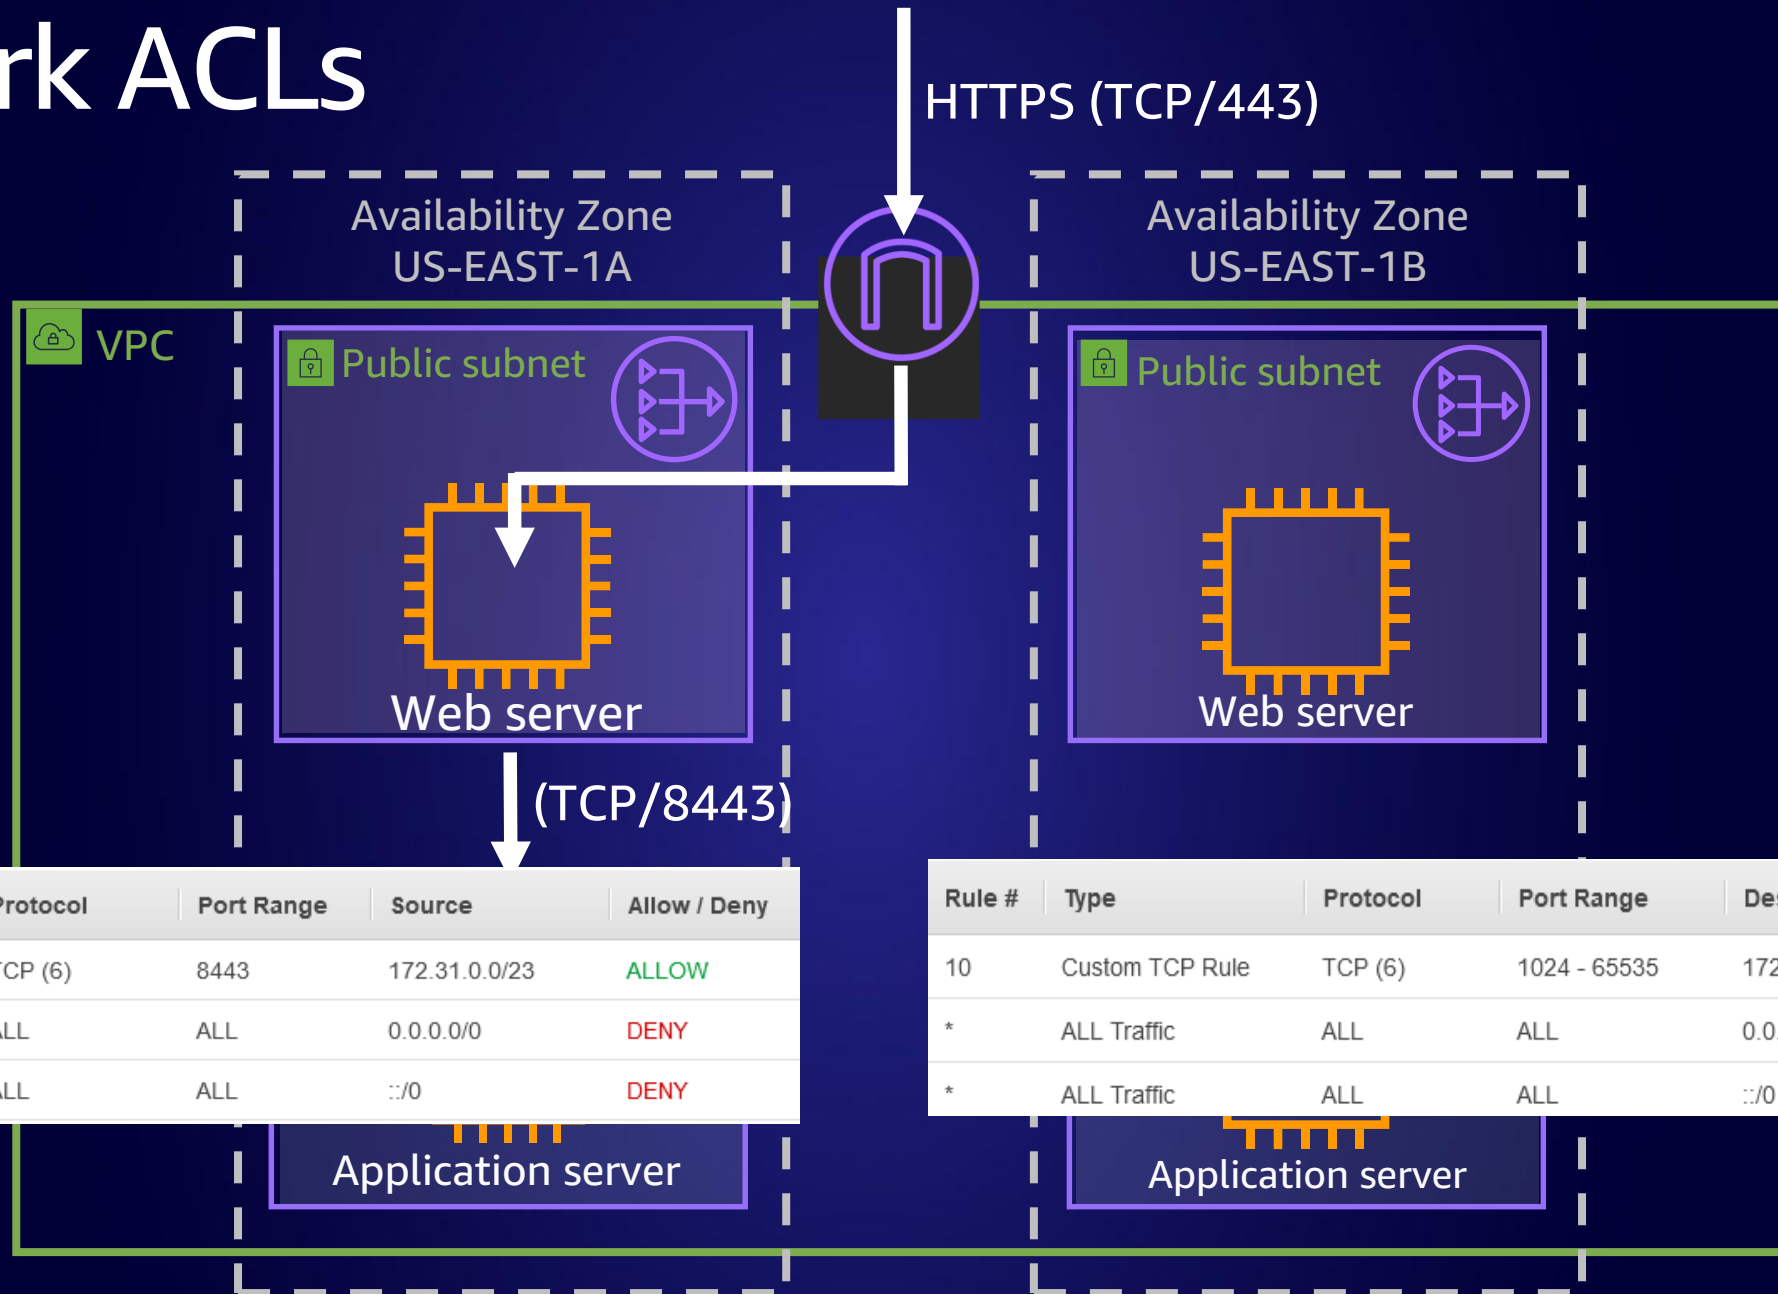
VPC Endpoint: [vpce-2222.foo.amazon.com](#)

Network Security


Network ACLs

- **Subnet** based security
- L3/L4 **Stateless** 
- Support ingress and egress rules
- **Both** “Allow” and “Deny” rules (order matters)
- Default: Allow All
- Limits:
 - Network ACLs per VPC: 200
 - Rules per network ACL: 20/40

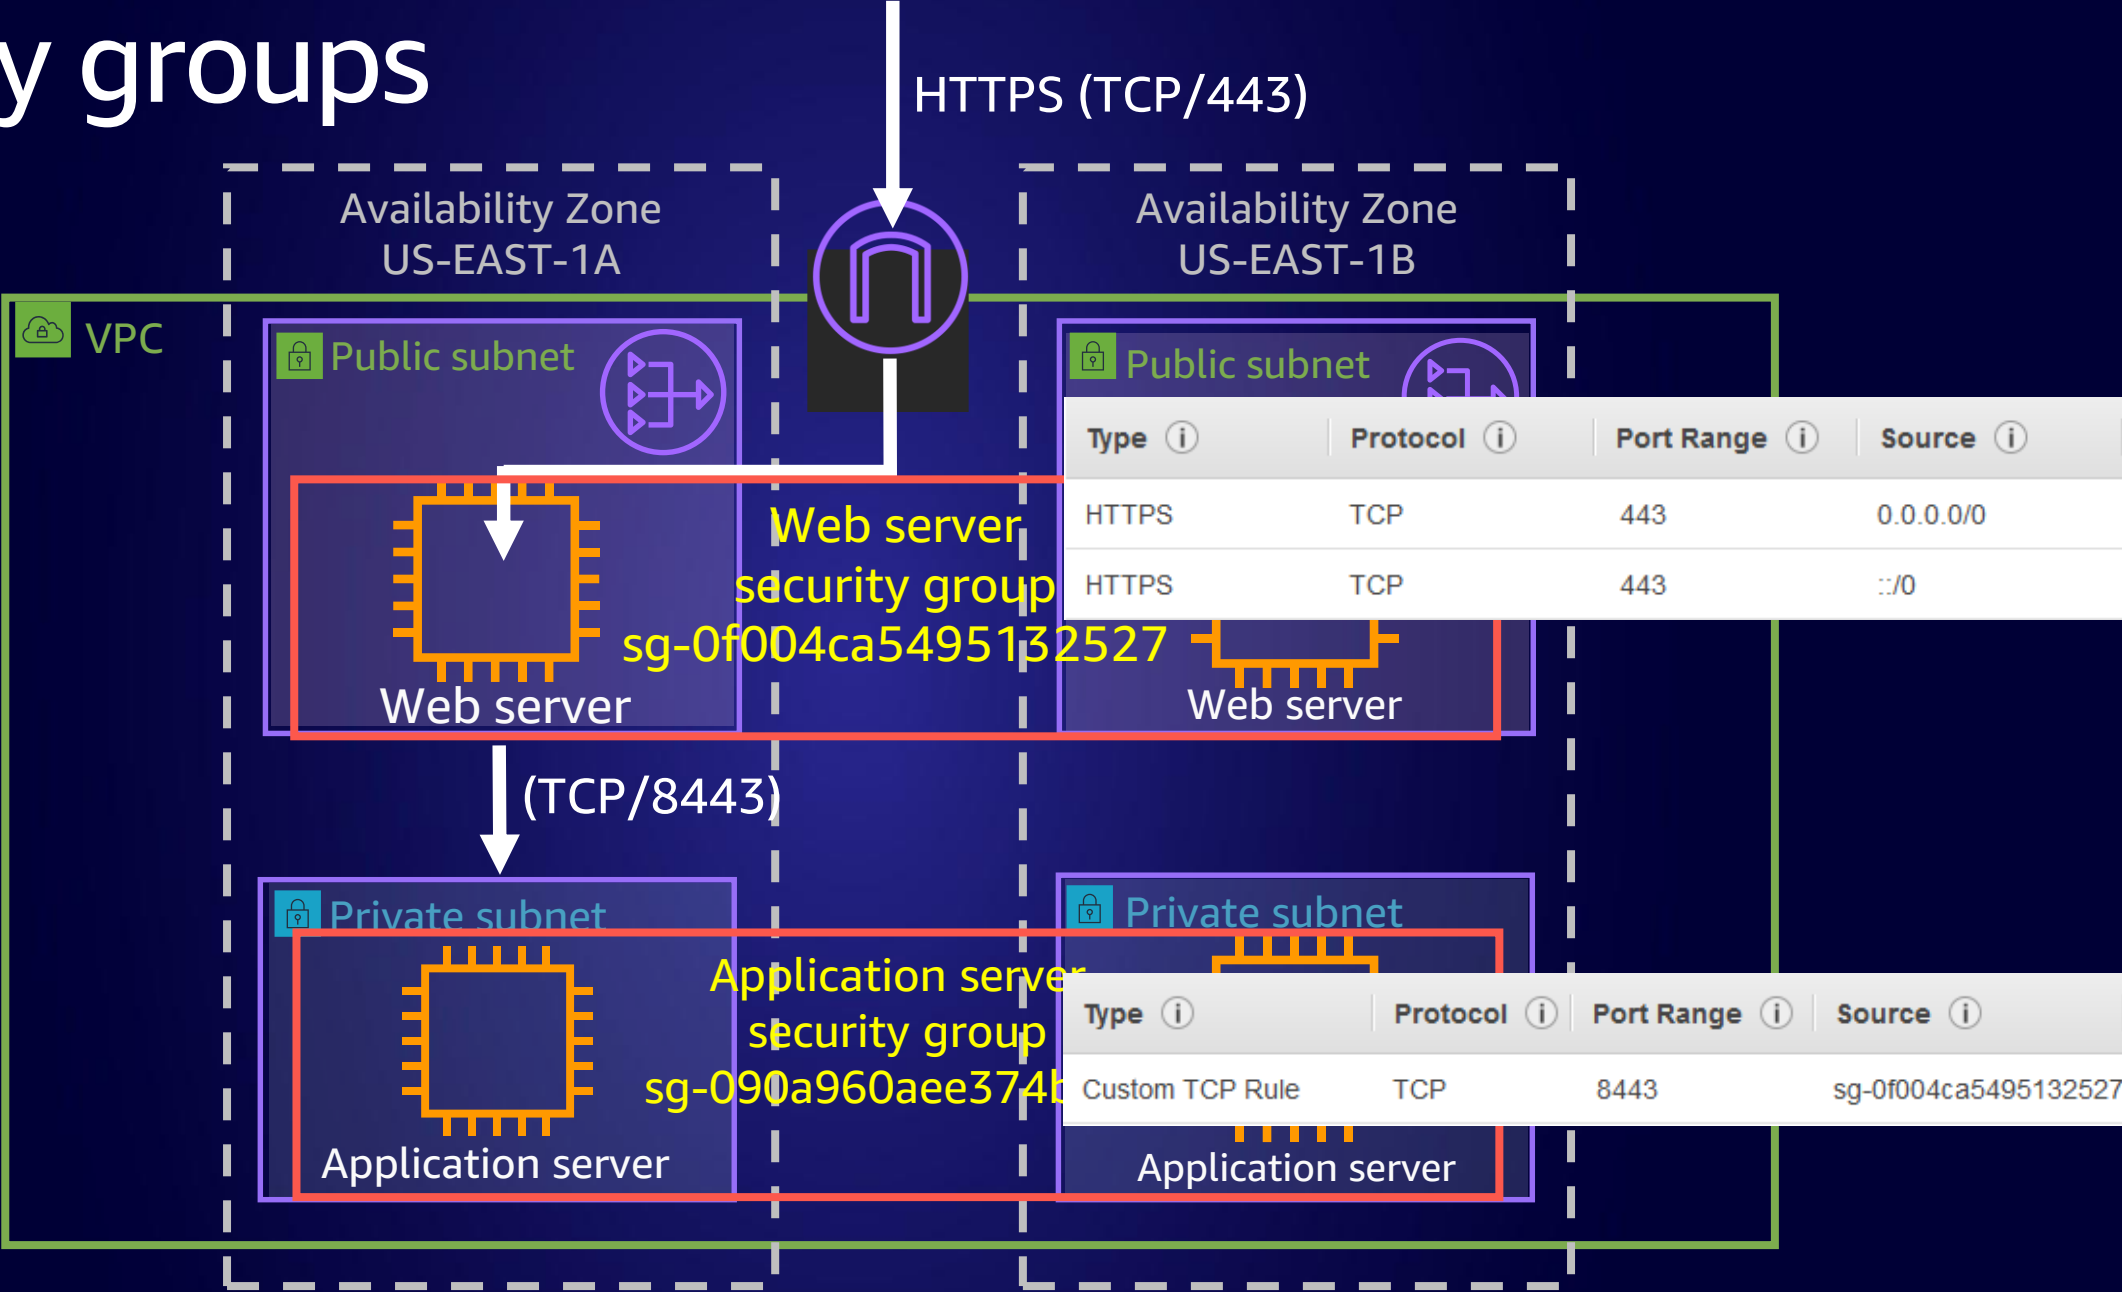
Network ACLs



Security groups

- **Instance**-level security
- Similar to host-based packet filter
- L3/L4 **Stateful** 
- **Only** “Allow” rules (order irrelevant)
- Support ingress and egress rules
- Allow reference of other security groups
 - Abstract function from IP addresses
- Limits:
 - Security groups per ENI: 5
 - Rule per security group: 60

Security groups

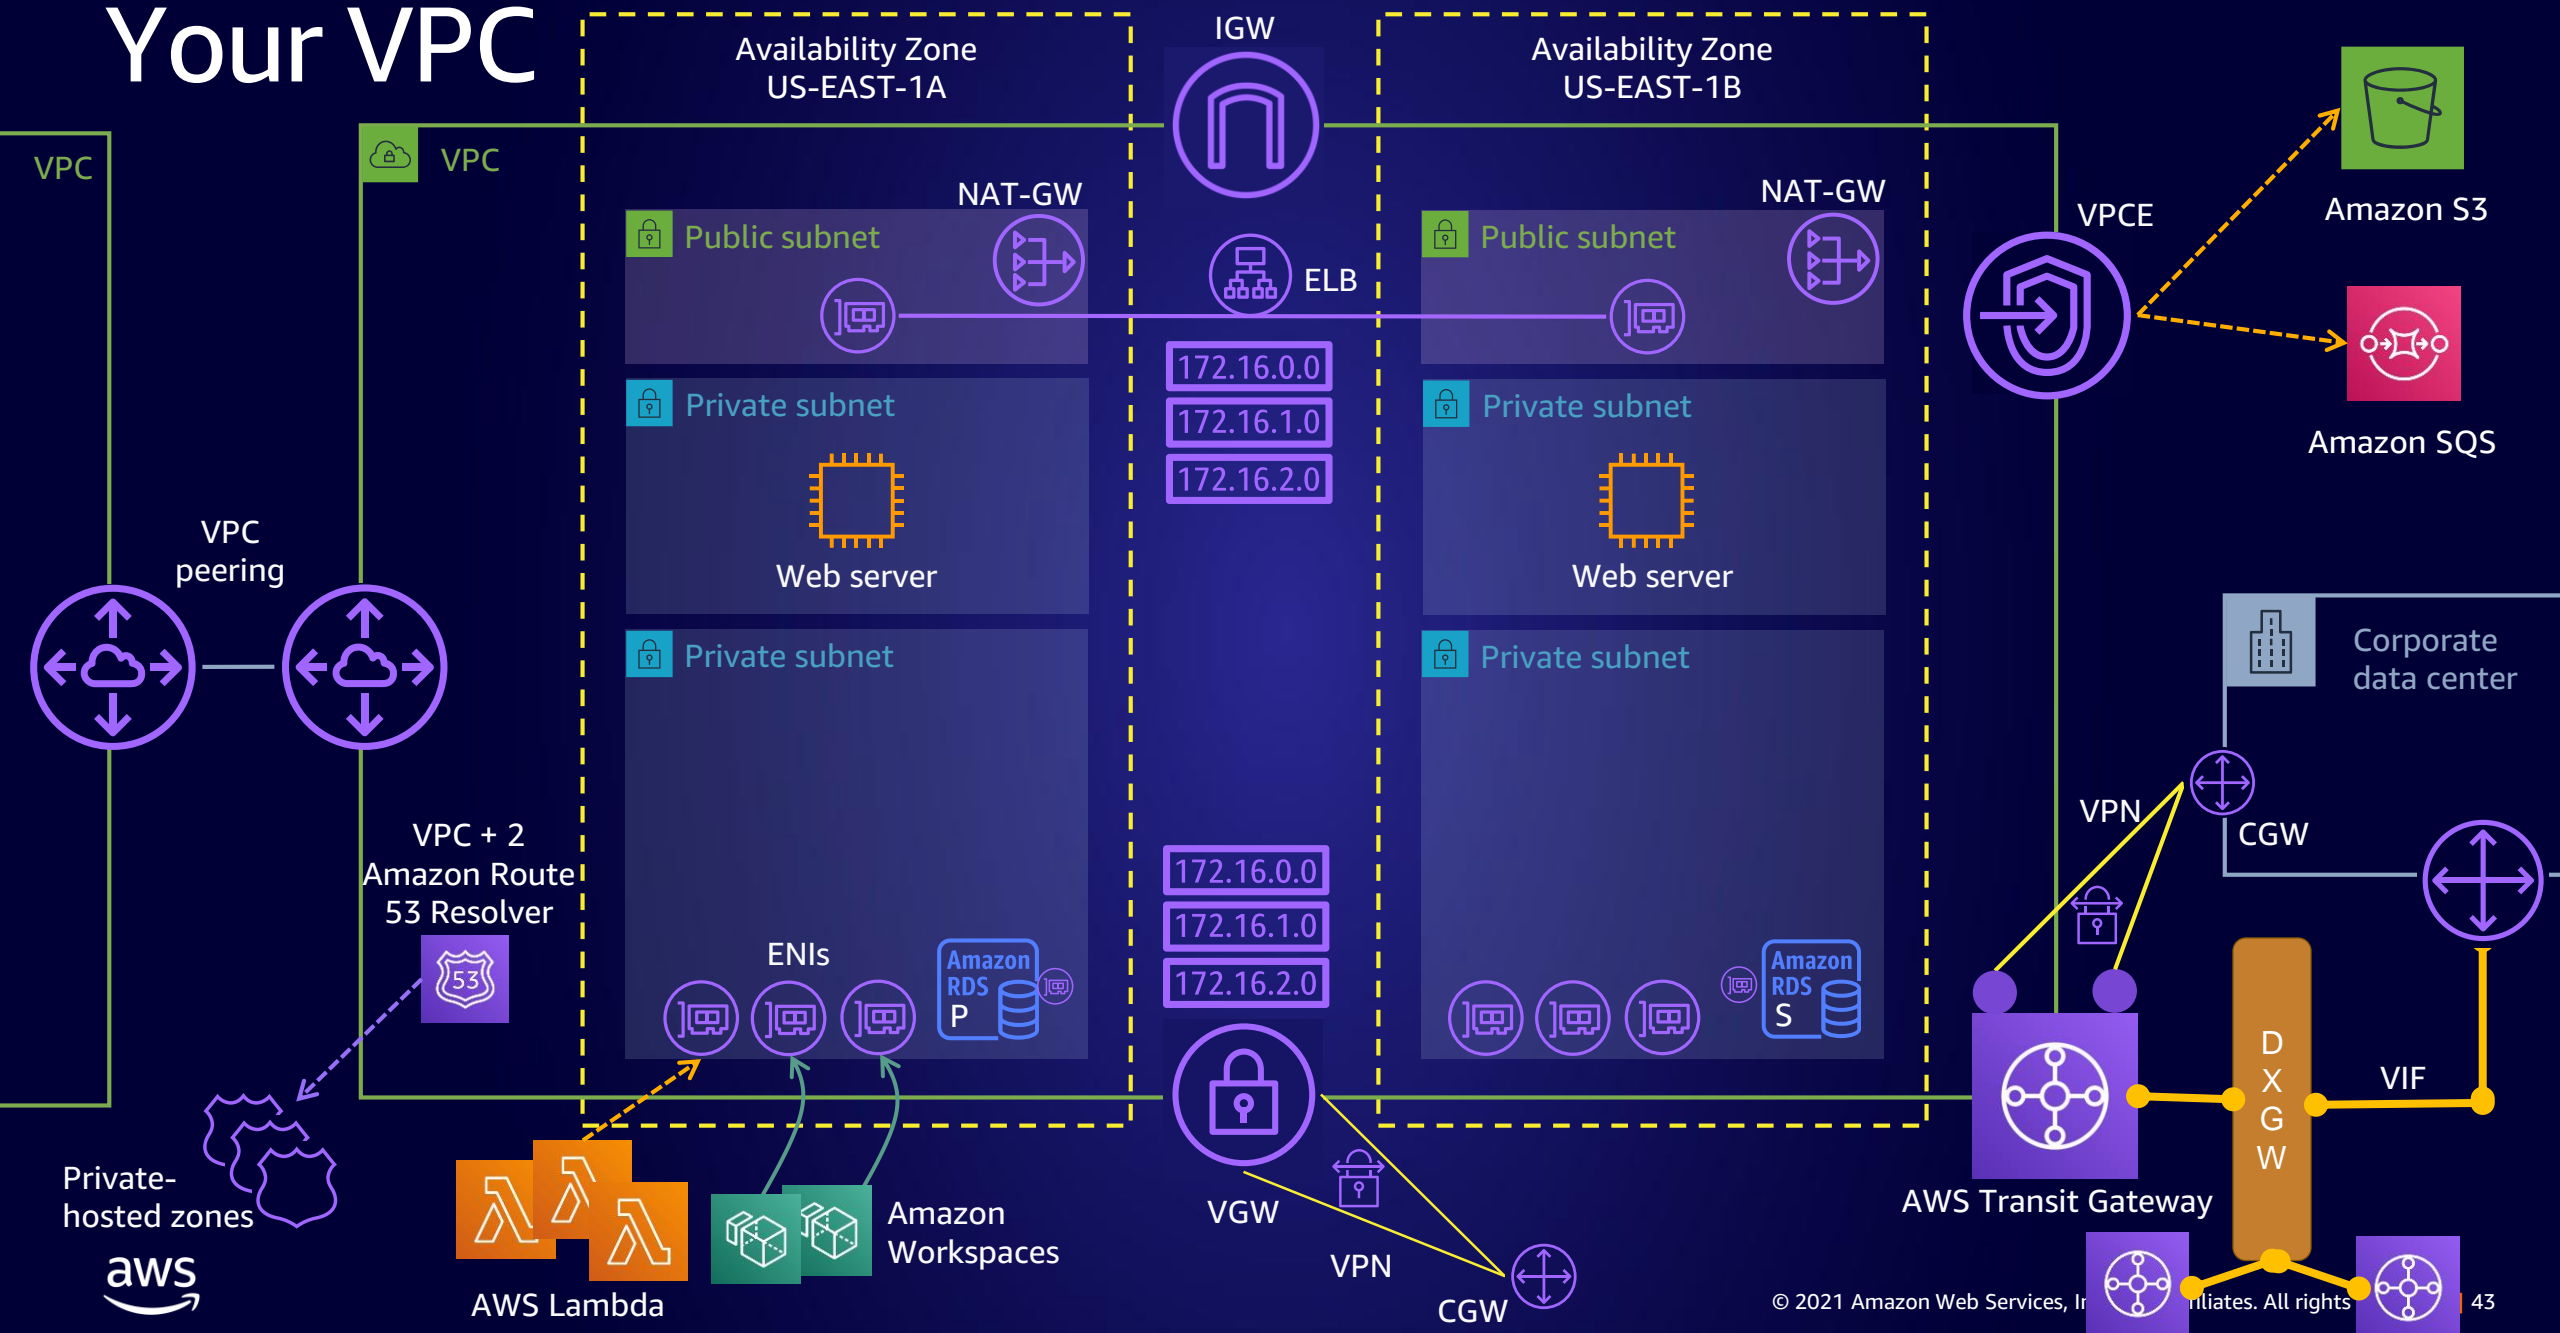


Security groups versus NACLs

Security group	Network ACL
Operates at instance level	Operates at subnet level
Supports allow rules only	Supports allow and deny rules
Is stateful: return traffic is automatically allowed regardless of any rules	Is stateless: return traffic must be explicitly allowed by rules
All rules evaluated before deciding whether to allow traffic	Rules evaluated in order when deciding whether to allow traffic
Applies only to instances explicitly associated with the security group	Automatically applies to all instances launched into associated subnets
Doesn't filter traffic to or from link-local addresses (169.254.0.0/16) or AWS-reserved IPv4 addresses; these are the first four IPv4 addresses of the subnet (including the Amazon VPC DNS server)	

Bringing It All Together

Your VPC




The famous from Amazon CTO Werner Vogels

“Everything fails, all the time”

Learn Networking with AWS Training and Certification

Resources created by the experts at AWS to help you build and validate networking skills



Learn online with 15+ free digital courses, including:
[Configuring and Deploying VPCs with Multiple Subnets](#) (1 hour)
and [Transit Gateway Networking and Scaling](#) (2 hours)



Build credibility and confidence with AWS Certification,
including:
[AWS Certified Advanced Networking – Specialty](#)

Visit aws.training/Networking

Thank you!

Shih-Yong Wang
shihyong@amazon.com

